

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from matplotlib import rcParams
import plotly.express as px
import numpy as np
import warnings
warnings = 'ignore'
%matplotlib inline
```

Dataset Description

link to dataset :

<https://www.kaggle.com/competitions/rossmann-store-sales/data>

You are provided with historical sales data for 1,115 Rossmann stores. The task is to forecast the "Sales" column for the test set. Note that some stores in the dataset were temporarily closed for refurbishment.

Files train.csv - historical data including Sales

test.csv - historical data excluding Sales

sample_submission.csv - a sample submission file in the correct format

store.csv - supplemental information about the stores

Data fields

Most of the fields are self-explanatory. The following are descriptions for those that aren't.

Id - an Id that represents a (Store, Date) duple within the test set

Store - a unique Id for each store

Sales - the turnover for any given day (this is what you are predicting)

Customers - the number of customers on a given day

Open - an indicator for whether the store was open: 0 = closed, 1 = open

StateHoliday - indicates a state holiday. Normally all stores, with few exceptions, are closed on state holidays. Note that all schools are closed on public holidays and weekends. a = public holiday, b = Easter holiday, c = Christmas, 0 = None

SchoolHoliday - indicates if the (Store, Date) was affected by the closure of public schools

StoreType - differentiates between 4 different store models: a, b, c, d

Assortment - describes an assortment level: a = basic, b = extra, c = extended

CompetitionDistance - distance in meters to the nearest competitor store

CompetitionOpenSince[Month/Year] - gives the approximate year and month of the time the nearest competitor was opened

Promo - indicates whether a store is running a promo on that day

Promo2 - Promo2 is a continuing and consecutive promotion for some stores: 0 = store is not participating, 1 = store is participating

Promo2Since[Year/Week] - describes the year and calendar week when the store started participating in Promo2

PromoInterval - describes the consecutive intervals Promo2 is started, naming the months the promotion is started anew. E.g. "Feb,May,Aug,Nov" means each round starts in February, May, August, November of any given year for that store

```
sns.set_style('darkgrid')
plt.rcParams['font.size'] = 14
plt.rcParams['figure.figsize'] = (10, 6)
plt.rcParams['figure.facecolor'] = '#00000000'
import warnings
warnings.filterwarnings("ignore")
c_green = '#6DF10C'
c_yellow = '#F5DD0D'
c_cyan = '#0FFDEF'
c_blue = '#0141DE'
c_blue_light = '#2775FD'
c_purple = '#FF0DE5'
c_green_dark = '#1BB200'
e = np.e

ross_df = pd.read_csv('csv\\train.csv')
store_df = pd.read_csv('csv\\store.csv')
test_df = pd.read_csv('csv\\test.csv')
```

ross_df

Promo	Store	DayOfWeek	Date	Sales	Customers	Open	
0	1	5	2015-07-31	5263	555	1	1
1	2	5	2015-07-31	6064	625	1	1
2	3	5	2015-07-31	8314	821	1	1
3	4	5	2015-07-31	13995	1498	1	1
4	5	5	2015-07-31	4822	559	1	1

...
1017204	1111	2	2013-01-01	0	0	0	0
1017205	1112	2	2013-01-01	0	0	0	0
1017206	1113	2	2013-01-01	0	0	0	0
1017207	1114	2	2013-01-01	0	0	0	0
1017208	1115	2	2013-01-01	0	0	0	0

	StateHoliday	SchoolHoliday
0	0	1
1	0	1
2	0	1
3	0	1
4	0	1
...
1017204	a	1
1017205	a	1
1017206	a	1
1017207	a	1
1017208	a	1

[1017209 rows x 9 columns]

store_df

	Store	StoreType	Assortment	CompetitionDistance	\
0	1	c	a	1270.0	
1	2	a	a	570.0	
2	3	a	a	14130.0	
3	4	c	c	620.0	
4	5	a	a	29910.0	
...	
1110	1111	a	a	1900.0	
1111	1112	c	c	1880.0	
1112	1113	a	c	9260.0	
1113	1114	a	c	870.0	
1114	1115	d	c	5350.0	

	CompetitionOpenSinceMonth	CompetitionOpenSinceYear	Promo2	\
0	9.0	2008.0	0	
1	11.0	2007.0	1	
2	12.0	2006.0	1	
3	9.0	2009.0	0	
4	4.0	2015.0	0	
...	

1110	6.0	2014.0	1
1111	4.0	2006.0	0
1112	NaN	NaN	0
1113	NaN	NaN	0
1114	NaN	NaN	1

	Promo2SinceWeek	Promo2SinceYear	PromoInterval
0	NaN	NaN	NaN
1	13.0	2010.0	Jan, Apr, Jul, Oct
2	14.0	2011.0	Jan, Apr, Jul, Oct
3	NaN	NaN	NaN
4	NaN	NaN	NaN
...
1110	31.0	2013.0	Jan, Apr, Jul, Oct
1111	NaN	NaN	NaN
1112	NaN	NaN	NaN
1113	NaN	NaN	NaN
1114	22.0	2012.0	Mar, Jun, Sept, Dec

[1115 rows x 10 columns]

```
merged_train_df = ross_df.merge(store_df, how='left' , on='Store')
merged_test_df = test_df.merge(store_df, how='left' , on='Store')
```

merged_train_df

Promo \	Store	DayOfWeek	Date	Sales	Customers	Open	
0	1	5	2015-07-31	5263	555	1	1
1	2	5	2015-07-31	6064	625	1	1
2	3	5	2015-07-31	8314	821	1	1
3	4	5	2015-07-31	13995	1498	1	1
4	5	5	2015-07-31	4822	559	1	1
...
1017204	1111	2	2013-01-01	0	0	0	0
1017205	1112	2	2013-01-01	0	0	0	0
1017206	1113	2	2013-01-01	0	0	0	0
1017207	1114	2	2013-01-01	0	0	0	0
1017208	1115	2	2013-01-01	0	0	0	0

CompetitionDistance	StateHoliday	SchoolHoliday	StoreType	Assortment
0	0	1	c	a
1270.0				
1	0	1	a	a
570.0				
2	0	1	a	a
14130.0				
3	0	1	c	c
620.0				
4	0	1	a	a
29910.0				
...
...				
1017204	a	1	a	a
1900.0				
1017205	a	1	c	c
1880.0				
1017206	a	1	a	c
9260.0				
1017207	a	1	a	c
870.0				
1017208	a	1	d	c
5350.0				

Promo2	CompetitionOpenSinceMonth	CompetitionOpenSinceYear	
0	9.0	2008.0	0
1	11.0	2007.0	1
2	12.0	2006.0	1
3	9.0	2009.0	0
4	4.0	2015.0	0
...
1017204	6.0	2014.0	1
1017205	4.0	2006.0	0
1017206	NaN	NaN	0
1017207	NaN	NaN	0
1017208	NaN	NaN	1

	Promo2SinceWeek	Promo2SinceYear	PromoInterval				
0	NaN	NaN	NaN				
1	13.0	2010.0	Jan, Apr, Jul, Oct				
2	14.0	2011.0	Jan, Apr, Jul, Oct				
3	NaN	NaN	NaN				
4	NaN	NaN	NaN				
...				
1017204	31.0	2013.0	Jan, Apr, Jul, Oct				
1017205	NaN	NaN	NaN				
1017206	NaN	NaN	NaN				
1017207	NaN	NaN	NaN				
1017208	22.0	2012.0	Mar, Jun, Sept, Dec				
[1017209 rows x 18 columns]							
merged_test_df							
	Id	Store	DayOfWeek	Date	Open	Promo	
StateHoliday \							
0	1	1	4	2015-09-17	1.0	1	0
1	2	3	4	2015-09-17	1.0	1	0
2	3	7	4	2015-09-17	1.0	1	0
3	4	8	4	2015-09-17	1.0	1	0
4	5	9	4	2015-09-17	1.0	1	0
...
41083	41084	1111	6	2015-08-01	1.0	0	0
41084	41085	1112	6	2015-08-01	1.0	0	0
41085	41086	1113	6	2015-08-01	1.0	0	0
41086	41087	1114	6	2015-08-01	1.0	0	0
41087	41088	1115	6	2015-08-01	1.0	0	0
	SchoolHoliday	StoreType	Assortment	CompetitionDistance	\		
0	0	c	a	1270.0			
1	0	a	a	14130.0			
2	0	a	c	24000.0			
3	0	a	a	7520.0			
4	0	a	c	2030.0			
...			
41083	0	a	a	1900.0			
41084	0	c	c	1880.0			

41085	0	a	c	9260.0
41086	0	a	c	870.0
41087	1	d	c	5350.0

	CompetitionOpenSinceMonth	CompetitionOpenSinceYear	Promo2	\
0	9.0	2008.0	0	
1	12.0	2006.0	1	
2	4.0	2013.0	0	
3	10.0	2014.0	0	
4	8.0	2000.0	0	
...
41083	6.0	2014.0	1	
41084	4.0	2006.0	0	
41085	NaN	NaN	0	
41086	NaN	NaN	0	
41087	NaN	NaN	1	

	Promo2SinceWeek	Promo2SinceYear	PromoInterval
0	NaN	NaN	NaN
1	14.0	2011.0	Jan, Apr, Jul, Oct
2	NaN	NaN	NaN
3	NaN	NaN	NaN
4	NaN	NaN	NaN
...
41083	31.0	2013.0	Jan, Apr, Jul, Oct
41084	NaN	NaN	NaN
41085	NaN	NaN	NaN
41086	NaN	NaN	NaN
41087	22.0	2012.0	Mar, Jun, Sept, Dec

[41088 rows x 17 columns]

Preprocessing and Feature Engineering

Let's take a look at the available columns, and figure out if we can create new columns or apply any useful transformations.

```
merged_train_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1017209 entries, 0 to 1017208
Data columns (total 18 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Store               1017209 non-null  int64
1   DayOfWeek           1017209 non-null  int64
2   Date               1017209 non-null  object
3   Sales              1017209 non-null  int64
4   Customers           1017209 non-null  int64
```

```

5   Open          1017209 non-null int64
6   Promo         1017209 non-null int64
7   StateHoliday  1017209 non-null object
8   SchoolHoliday 1017209 non-null int64
9   StoreType     1017209 non-null object
10  Assortment     1017209 non-null object
11  CompetitionDistance 1014567 non-null float64
12  CompetitionOpenSinceMonth 693861 non-null float64
13  CompetitionOpenSinceYear 693861 non-null float64
14  Promo2        1017209 non-null int64
15  Promo2SinceWeek 509178 non-null float64
16  Promo2SinceYear 509178 non-null float64
17  PromoInterval 509178 non-null object

```

dtypes: float64(5), int64(8), object(5)

memory usage: 139.7+ MB

```

def split_date(df):
    df['Date'] = pd.to_datetime(df['Date'])
    df['Year'] = df.Date.dt.year
    df['Month'] = df.Date.dt.month
    df['Day'] = df.Date.dt.day
    df['WeekOfYear'] = df.Date.dt.isocalendar().week

```

split_date(merged_train_df)

split_date(merged_test_df)

merged_train_df

	Store	DayOfWeek	Date	Sales	Customers	Open	Promo	\
0	1	5	2015-07-31	5263	555	1	1	
1	2	5	2015-07-31	6064	625	1	1	
2	3	5	2015-07-31	8314	821	1	1	
3	4	5	2015-07-31	13995	1498	1	1	
4	5	5	2015-07-31	4822	559	1	1	
...	
1017204	1111	2	2013-01-01	0	0	0	0	
1017205	1112	2	2013-01-01	0	0	0	0	
1017206	1113	2	2013-01-01	0	0	0	0	
1017207	1114	2	2013-01-01	0	0	0	0	
1017208	1115	2	2013-01-01	0	0	0	0	

	StateHoliday	SchoolHoliday	StoreType	...
CompetitionOpenSinceMonth				
0	0	1	c	...
9.0				
1	0	1	a	...
11.0				
2	0	1	a	...
12.0				
3	0	1	c	...

9.0				
4	0	1	a	...
4.0				
...
...				
1017204	a	1	a	...
6.0				
1017205	a	1	c	...
4.0				
1017206	a	1	a	...
NaN				
1017207	a	1	a	...
NaN				
1017208	a	1	d	...
NaN				

	CompetitionOpenSinceYear	Promo2	Promo2SinceWeek
Promo2SinceYear \			
0	2008.0	0	NaN
NaN			
1	2007.0	1	13.0
2010.0			
2	2006.0	1	14.0
2011.0			
3	2009.0	0	NaN
NaN			
4	2015.0	0	NaN
NaN			
...
...			
1017204	2014.0	1	31.0
2013.0			
1017205	2006.0	0	NaN
NaN			
1017206	NaN	0	NaN
NaN			
1017207	NaN	0	NaN
NaN			
1017208	NaN	1	22.0
2012.0			

	PromoInterval	Year	Month	Day	WeekOfYear
0	NaN	2015	7	31	31
1	Jan, Apr, Jul, Oct	2015	7	31	31
2	Jan, Apr, Jul, Oct	2015	7	31	31
3	NaN	2015	7	31	31
4	NaN	2015	7	31	31
...
1017204	Jan, Apr, Jul, Oct	2013	1	1	1

1017205		NaN	2013	1	1	1
1017206		NaN	2013	1	1	1
1017207		NaN	2013	1	1	1
1017208	Mar,Jun,Sept,Dec		2013	1	1	1

[1017209 rows x 22 columns]

```
merged_train_df[merged_train_df.Open == 0].Sales.value_counts()
```

Sales

0 172817

Name: count, dtype: int64

```
merged_train_df = merged_train_df[merged_train_df.Open == 1].copy()
```

merged_train_df

	Store	DayOfWeek	Date	Sales	Customers	Open	Promo	\
0	1	5	2015-07-31	5263	555	1	1	
1	2	5	2015-07-31	6064	625	1	1	
2	3	5	2015-07-31	8314	821	1	1	
3	4	5	2015-07-31	13995	1498	1	1	
4	5	5	2015-07-31	4822	559	1	1	
...
1016776	682	2	2013-01-01	3375	566	1	0	
1016827	733	2	2013-01-01	10765	2377	1	0	
1016863	769	2	2013-01-01	5035	1248	1	0	
1017042	948	2	2013-01-01	4491	1039	1	0	
1017190	1097	2	2013-01-01	5961	1405	1	0	

	StateHoliday	SchoolHoliday	StoreType	...
CompetitionOpenSinceMonth				\
0	0	1	c	...
9.0				
1	0	1	a	...
11.0				
2	0	1	a	...
12.0				
3	0	1	c	...
9.0				
4	0	1	a	...
4.0				
...
...				
1016776	a	1	b	...
9.0				
1016827	a	1	b	...
10.0				
1016863	a	1	b	...
NaN				
1017042	a	1	b	...

```

NaN
1017190          a          1          b  ...
3.0

      CompetitionOpenSinceYear  Promo2  Promo2SinceWeek
Promo2SinceYear  \
0                2008.0          0                NaN
NaN
1                2007.0          1                13.0
2010.0
2                2006.0          1                14.0
2011.0
3                2009.0          0                NaN
NaN
4                2015.0          0                NaN
NaN
...                ...          ...                ...
...
1016776          2006.0          0                NaN
NaN
1016827          1999.0          0                NaN
NaN
1016863          NaN           1                48.0
2012.0
1017042          NaN           0                NaN
NaN
1017190          2002.0          0                NaN
NaN

      PromoInterval  Year  Month  Day  WeekOfYear
0                NaN  2015     7   31           31
1      Jan, Apr, Jul, Oct  2015     7   31           31
2      Jan, Apr, Jul, Oct  2015     7   31           31
3                NaN  2015     7   31           31
4                NaN  2015     7   31           31
...                ...    ...    ...    ...
1016776          NaN  2013     1    1            1
1016827          NaN  2013     1    1            1
1016863  Jan, Apr, Jul, Oct  2013     1    1            1
1017042          NaN  2013     1    1            1
1017190          NaN  2013     1    1            1

[844392 rows x 22 columns]

```

Competition

Next, we can use the columns `CompetitionOpenSince[Month/Year]` columns from `store_df` to compute the number of months for which a competitor has been open near the store.

```
def comp_months(df):
    df['CompetitionOpen'] = 12 * (df.Year -
df.CompetitionOpenSinceYear) + (df.Month -
df.CompetitionOpenSinceMonth)
    df['CompetitionOpen'] = df['CompetitionOpen'].map(lambda x: 0 if x
< 0 else x).fillna(0)
```

```
comp_months(merged_train_df)
```

```
comp_months(merged_test_df)
```

```
merged_train_df
```

	Store	DayOfWeek	Date	Sales	Customers	Open	Promo	\
0	1	5	2015-07-31	5263	555	1	1	
1	2	5	2015-07-31	6064	625	1	1	
2	3	5	2015-07-31	8314	821	1	1	
3	4	5	2015-07-31	13995	1498	1	1	
4	5	5	2015-07-31	4822	559	1	1	
...
1016776	682	2	2013-01-01	3375	566	1	0	
1016827	733	2	2013-01-01	10765	2377	1	0	
1016863	769	2	2013-01-01	5035	1248	1	0	
1017042	948	2	2013-01-01	4491	1039	1	0	
1017190	1097	2	2013-01-01	5961	1405	1	0	

	StateHoliday	SchoolHoliday	StoreType	...
CompetitionOpenSinceYear				\
0	0	1	c	...
2008.0				
1	0	1	a	...
2007.0				
2	0	1	a	...
2006.0				
3	0	1	c	...
2009.0				
4	0	1	a	...
2015.0				
...
...				
1016776	a	1	b	...
2006.0				
1016827	a	1	b	...
1999.0				
1016863	a	1	b	...
NaN				
1017042	a	1	b	...
NaN				
1017190	a	1	b	...
2002.0				

Year \	Promo2	Promo2SinceWeek	Promo2SinceYear	PromoInterval
0	0	NaN	NaN	NaN
2015				
1	1	13.0	2010.0	Jan, Apr, Jul, Oct
2015				
2	1	14.0	2011.0	Jan, Apr, Jul, Oct
2015				
3	0	NaN	NaN	NaN
2015				
4	0	NaN	NaN	NaN
2015				
...
..				
1016776	0	NaN	NaN	NaN
2013				
1016827	0	NaN	NaN	NaN
2013				
1016863	1	48.0	2012.0	Jan, Apr, Jul, Oct
2013				
1017042	0	NaN	NaN	NaN
2013				
1017190	0	NaN	NaN	NaN
2013				

	Month	Day	WeekOfYear	CompetitionOpen
0	7	31	31	82.0
1	7	31	31	92.0
2	7	31	31	103.0
3	7	31	31	70.0
4	7	31	31	3.0
...
1016776	1	1	1	76.0
1016827	1	1	1	159.0
1016863	1	1	1	0.0
1017042	1	1	1	0.0
1017190	1	1	1	130.0

[844392 rows x 23 columns]

merged_train_df.columns

Index(['Store', 'DayOfWeek', 'Date', 'Sales', 'Customers', 'Open', 'Promo',
'StateHoliday', 'SchoolHoliday', 'StoreType', 'Assortment',
'CompetitionDistance', 'CompetitionOpenSinceMonth',
'CompetitionOpenSinceYear', 'Promo2', 'Promo2SinceWeek',
'Promo2SinceYear', 'PromoInterval', 'Year', 'Month', 'Day',
'WeekOfYear', 'CompetitionOpen'],
dtype='object')

```

input_cols = ['Store' , 'DayOfWeek' , 'Promo' , 'StateHoliday' ,
'Assortment' ,
'CompetitionDistance' , 'Day' , 'Month' , 'Year' ,
'WeekOfYear' ,
'Promo2' , 'CompetitionOpen' , 'SchoolHoliday' ,
'StoreType' ]
target_col = 'Sales'

```

```

train_inputs = merged_train_df[input_cols].copy()
train_targets = merged_train_df[target_col].copy()

```

```

test_inputs = merged_test_df[input_cols].copy()

```

```

train_inputs

```

	Store	DayOfWeek	Promo	StateHoliday	Assortment
CompetitionDistance \					
0	1	5	1	0	a
1270.0					
1	2	5	1	0	a
570.0					
2	3	5	1	0	a
14130.0					
3	4	5	1	0	c
620.0					
4	5	5	1	0	a
29910.0					
...
...					
1016776	682	2	0	a	a
150.0					
1016827	733	2	0	a	b
860.0					
1016863	769	2	0	a	b
840.0					
1017042	948	2	0	a	b
1430.0					
1017190	1097	2	0	a	b
720.0					

	Day	Month	Year	WeekOfYear	Promo2	CompetitionOpen
SchoolHoliday \						
0	31	7	2015	31	0	82.0
1						
1	31	7	2015	31	1	92.0
1						
2	31	7	2015	31	1	103.0
1						
3	31	7	2015	31	0	70.0
1						

4	31	7	2015	31	0	3.0
1						
...
...						
1016776	1	1	2013	1	0	76.0
1						
1016827	1	1	2013	1	0	159.0
1						
1016863	1	1	2013	1	1	0.0
1						
1017042	1	1	2013	1	0	0.0
1						
1017190	1	1	2013	1	0	130.0
1						

	StoreType
0	c
1	a
2	a
3	c
4	a
...	...
1016776	b
1016827	b
1016863	b
1017042	b
1017190	b

[844392 rows x 14 columns]

train_inputs.columns

```
Index(['Store', 'DayOfWeek', 'Promo', 'StateHoliday', 'Assortment',
      'CompetitionDistance', 'Day', 'Month', 'Year', 'WeekOfYear',
      'Promo2',
      'CompetitionOpen', 'SchoolHoliday', 'StoreType'],
      dtype='object')
```

```
numeric_cols = ['Store' , 'Promo' , 'SchoolHoliday' ,
                'CompetitionDistance' , 'CompetitionOpen' , 'Promo2' ,
                'Day' , 'Month' , 'Year' , 'WeekOfYear']
```

```
categorical_cols = [ 'DayOfWeek' , 'StoreType' , 'Assortment']
```

```
train_inputs[numeric_cols].isna().sum().sort_values(ascending=False)
```

CompetitionDistance	2186
Store	0
Promo	0
SchoolHoliday	0
CompetitionOpen	0

```
Promo2      0
Day          0
Month        0
Year         0
WeekOfYear   0
dtype: int64

test_inputs[numeric_cols].isna().sum().sort_values(ascending=False)

CompetitionDistance    96
Store                  0
Promo                  0
SchoolHoliday          0
CompetitionOpen        0
Promo2                 0
Day                    0
Month                  0
Year                   0
WeekOfYear             0
dtype: int64
```

Seems like competition distance is the only missing value, and we can simply fill it with the highest value (to indicate that competition is very far away).

```
max_distance = train_inputs.CompetitionDistance.max()

train_inputs['CompetitionDistance'].fillna(max_distance, inplace=True)
test_inputs['CompetitionDistance'].fillna(max_distance, inplace=True)

train_inputs[numeric_cols].isna().sum().sort_values(ascending=False)

Store          0
Promo          0
SchoolHoliday  0
CompetitionDistance  0
CompetitionOpen  0
Promo2         0
Day            0
Month          0
Year           0
WeekOfYear     0
dtype: int64

train_inputs
```

	Store	DayOfWeek	Promo	StateHoliday	Assortment
CompetitionDistance \					
0	1	5	1	0	a
1270.0					
1	2	5	1	0	a

570.0					
2	3	5	1	0	a
14130.0					
3	4	5	1	0	c
620.0					
4	5	5	1	0	a
29910.0					
...
...					
1016776	682	2	0	a	a
150.0					
1016827	733	2	0	a	b
860.0					
1016863	769	2	0	a	b
840.0					
1017042	948	2	0	a	b
1430.0					
1017190	1097	2	0	a	b
720.0					

	Day	Month	Year	Week0fYear	Promo2	CompetitionOpen
SchoolHoliday \						
0	31	7	2015	31	0	82.0
1						
1	31	7	2015	31	1	92.0
1						
2	31	7	2015	31	1	103.0
1						
3	31	7	2015	31	0	70.0
1						
4	31	7	2015	31	0	3.0
1						
...
...						
1016776	1	1	2013	1	0	76.0
1						
1016827	1	1	2013	1	0	159.0
1						
1016863	1	1	2013	1	1	0.0
1						
1017042	1	1	2013	1	0	0.0
1						
1017190	1	1	2013	1	0	130.0
1						

	StoreType
0	c
1	a
2	a

```

3          c
4          a
...      ...
1016776    b
1016827    b
1016863    b
1017042    b
1017190    b

```

[844392 rows x 14 columns]

```

from sklearn.preprocessing import MinMaxScaler , MaxAbsScaler ,
StandardScaler
from sklearn.preprocessing import OneHotEncoder , LabelEncoder ,
TargetEncoder
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import SGDRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.tree import DecisionTreeRegressor
from xgboost import XGBRegressor
from lightgbm import LGBMRegressor
from sklearn.metrics import accuracy_score,mean_squared_error
import joblib

```

```

def rmse(targets , predictions):
    return np.sqrt(np.mean(np.square(targets-predictions)))

```

```

def mse(targets , predictions):
    return np.mean(np.square(targets-predictions))

```

train_inputs

	Store	DayOfWeek	Promo	StateHoliday	Assortment
CompetitionDistance \					
0	1	5	1	0	a
1270.0					
1	2	5	1	0	a
570.0					
2	3	5	1	0	a
14130.0					
3	4	5	1	0	c
620.0					
4	5	5	1	0	a
29910.0					
...
...					
1016776	682	2	0	a	a
150.0					

1016827	733	2	0	a	b
860.0					
1016863	769	2	0	a	b
840.0					
1017042	948	2	0	a	b
1430.0					
1017190	1097	2	0	a	b
720.0					

	Day	Month	Year	WeekOfYear	Promo2	CompetitionOpen
SchoolHoliday \						
0	31	7	2015	31	0	82.0
1						
1	31	7	2015	31	1	92.0
1						
2	31	7	2015	31	1	103.0
1						
3	31	7	2015	31	0	70.0
1						
4	31	7	2015	31	0	3.0
1						
...
...						
1016776	1	1	2013	1	0	76.0
1						
1016827	1	1	2013	1	0	159.0
1						
1016863	1	1	2013	1	1	0.0
1						
1017042	1	1	2013	1	0	0.0
1						
1017190	1	1	2013	1	0	130.0
1						

	StoreType
0	c
1	a
2	a
3	c
4	a
...	...
1016776	b
1016827	b
1016863	b
1017042	b
1017190	b

[844392 rows x 14 columns]

```

class RossmanSalesPrediction:

    def __init__(self, train_inputs, train_target, test_inputs,
numeric_cols, categorical_cols):
        self.train_inputs = train_inputs
        self.train_target = train_target
        self.test_inputs = test_inputs
        self.numeric_cols = numeric_cols
        self.categorical_cols = categorical_cols

    def impute_missing_features(self, imputer, **params):
        try:
            self.imputer =
imputer(**params).fit(self.train_inputs[self.numeric_cols])
            self.imputer_stats = list(self.imputer.statistics_)
            self.train_inputs[self.numeric_cols] =
self.imputer.transform(self.train_inputs[self.numeric_cols])
        except:
            raise ValueError('No Missing Data in train_inputs')

    ...

    Takes Input train_inputs and test_inputs and scales them according
to the scaler is provided,
    Function returns a list which contains scaled train_inputs and
test_inputs [train_inputs , test_inputs]
    ...

    def scale_num_features(self, scaler):
        self.scaler =
scaler().fit(self.train_inputs[self.numeric_cols])
        self.train_inputs[self.numeric_cols] =
self.scaler.transform(self.train_inputs[self.numeric_cols])
        self.test_inputs[self.numeric_cols] =
self.scaler.transform(self.test_inputs[self.numeric_cols])
        return [self.train_inputs, self.test_inputs]

    ...

    Takes input: (train_inputs and test_inputs) then encodes them
according to the Encoder is provided,
    Function returns a list which contains encoded train_inputs,
test_inputs, [X_train, X_test]
    ...

    def encode_cat_features(self, encoder, **params):
        self.encoder =
encoder(sparse=False ,handle_unknown='ignore').fit(self.train_inputs[s
elf.categorical_cols])
        self.encoded_cols =
list(self.encoder.get_feature_names_out(self.categorical_cols))
        self.train_inputs[self.encoded_cols] =

```

```

self.encoder.transform(self.train_inputs[self.categorical_cols])
    self.test_inputs[self.encoded_cols] =
self.encoder.transform(self.test_inputs[self.categorical_cols])
    self.X_train = self.train_inputs[self.numeric_cols +
self.encoded_cols]
    self.X_test = self.test_inputs[self.numeric_cols +
self.encoded_cols]
    return [self.X_train, self.X_test , self.encoded_cols]

...
Takes input: (split_size) then splits X_train according to the
split_size into X_val and val_target
Function returns a list which contains
[self.X_train,self.X_val,self.train_target,self.val_target]
'''

def split_df(self,split_size):
    self.X_train, self.X_val, self.train_target, self.val_target =
train_test_split(self.X_train, self.train_target, test_size=split_size
, random_state=42)
    return
[self.X_train,self.X_val,self.train_target,self.val_target]

def train_model(self,model,**params):
    self.model =
model(**params).fit(self.X_train,self.train_target)
    self.train_preds = self.model.predict(self.X_train)
    self.val_preds = self.model.predict(self.X_val)
    return [self.train_preds, self.val_preds, self.model]

def get_accuracy_scores(self):
    self.train_rmse = rmse(self.train_target,self.train_preds)
    self.val_rmse = rmse(self.val_target,self.val_preds)
    self.r_2_score =
self.model.score(self.X_train,self.train_target)
    self.accuracy =
self.model.score(self.X_train,self.train_target)*100
    return {
        'Model': self.model,
        'train_rmse' : self.train_rmse,
        'val_rmse' : self.val_rmse,
        'r_2_score' : self.r_2_score,
        'Model Accuracy' : self.accuracy
    }

def predict_input(self,new_input):
    self.input_df = pd.DataFrame([new_input])
    self.input_df[self.numeric_cols] =
self.imputer.transform(self.input_df[self.numeric_cols])

```

```

        self.input_df[self.numeric_cols] =
self.scaler.transform(self.input_df[self.numeric_cols])
        self.input_df[self.encoded_cols] =
self.encoder.transform(self.input_df[self.categorical_cols])
        X_input = self.input_df[self.numeric_cols + self.encoded_cols]
        prediction = self.model.predict(X_input)[0]
        return f'Prediction of {self.model} : {prediction} '

def get_importance_df(self):
    self.importance_df = pd.DataFrame({
        'feature': self.X_train.columns,
        'importance': self.model.feature_importances_
    }).sort_values('importance', ascending=False)
    return self.importance_df

def plot_importance(self):
    plt.figure(figsize=(10,10))
    plt.title('Feature Importance of model')
    sns.barplot(data=self.importance_df, x='importance',
y='feature')
    plt.show()

def import_model(self):
    self.rossman_sales_model = {
        'model': self.model,
        'imputer': self.imputer,
        'scaler': self.scaler,
        'encoder': self.encoder,
        'numeric_cols': self.numeric_cols,
        'categorical_cols': self.categorical_cols,
        'encoded_cols': self.encoded_cols,
        'X_train': self.X_train,
        'X_val': self.X_val,
        'X_test': self.X_test,
        'train_inputs': self.train_inputs,
        'test_inputs': self.test_inputs,
        'train_target': self.train_target,
        'val_target': self.val_target,
        'imputer_statistics': self.imputer_stats,
    }
    joblib.dump(self.rossman_sales_model, "files\\
rossman_sales_model.joblib")

def list_params(self):
    return ['train_inputs',
        'test_inputs',
        'train_target' ,
        'val_target' ,
        'numeric_cols',
        'categorical_cols',

```

```

        'encoded_cols',
        'imputer_statistics',
        'X_train',
        'X_val',
        'X_test',
        'encoder',
        'scaler',
        'Model',
        'train_rmse',
        'val_rmse',
        'r_2_score',
        'Model Accuracy']

def get_params(self):
    return {
        'train_inputs' : self.train_inputs,
        'test_inputs' : self.test_inputs,
        'train_target' : self.train_target,
        'val_target' : self.val_target,
        'numeric_cols' : self.numeric_cols,
        'categorical_cols' : self.categorical_cols,
        'encoded_cols' : self.encoded_cols,
        'imputer_statistics' : self.imputer_stats,
        'X_train' : self.X_train,
        'X_val' : self.X_val,
        'X_test' : self.X_test,
        'encoder' : self.encoder,
        'scaler' : self.scaler,
        'Model' : self.model,
        'train_rmse' : self.train_rmse,
        'val_rmse' : self.val_rmse,
        'r_2_score' : self.r_2_score,
        'Model Accuracy' : self.accuracy
    }

modell =
RossmanSalesPrediction(train_inputs,train_targets,test_inputs,numeric_
cols,categorical_cols)

imputer_params ={
    'strategy':'mean',
    'missing_values':np.nan
}

imputer =
modell.impute_missing_features(SimpleImputer,**imputer_params)
imputer

scaler = modell.scale_num_features(MinMaxScaler)

```

```

encoder = model1.encode_cat_features(OneHotEncoder)
splited = model1.split_df(0.1)
xgb_params = {
    'n_jobs': -1,
    'random_state': 42,
    'n_estimators': 1000,
    'learning_rate': 0.1,
    'max_depth': 10,
    'subsample': 0.9,
    'colsample_bytree': 0.7
}

from xgboost import XGBRegressor

xgb = model1.train_model(XGBRegressor, **xgb_params)
xgb

[array([11869.245 , 3592.7712, 8157.4316, ..., 6393.0723, 8789.75
,
        11786.23  ], dtype=float32),
 array([4679.055 , 8167.9214, 7206.459 , ..., 5404.9844, 4775.9844,
        7083.6616], dtype=float32),
 XGBRegressor(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=0.7, device=None,
early_stopping_rounds=None,
               enable_categorical=False, eval_metric=None,
feature_types=None,
               gamma=None, grow_policy=None, importance_type=None,
               interaction_constraints=None, learning_rate=0.1,
max_bin=None,
               max_cat_threshold=None, max_cat_to_onehot=None,
               max_delta_step=None, max_depth=10, max_leaves=None,
               min_child_weight=None, missing=nan,
monotone_constraints=None,
               multi_strategy=None, n_estimators=1000, n_jobs=-1,
               num_parallel_tree=None, random_state=42, ...)]

xgb_scores = model1.get_accuracy_scores()
xgb_scores

{'Model': XGBRegressor(base_score=None, booster=None, callbacks=None,
                       colsample_bylevel=None, colsample_bynode=None,
                       colsample_bytree=0.7, device=None,
early_stopping_rounds=None,
                       enable_categorical=False, eval_metric=None,
feature_types=None,
                       gamma=None, grow_policy=None, importance_type=None,
                       interaction_constraints=None, learning_rate=0.1,

```



```

max_bin=None,
        max_cat_threshold=None, max_cat_to_onehot=None,
        max_delta_step=None, max_depth=10, max_leaves=None,
        min_child_weight=None, missing=nan,
monotone_constraints=None,
        multi_strategy=None, n_estimators=1000, n_jobs=-1,
        num_parallel_tree=None, random_state=42, ...),
'train_rmse': 476.5882506580303,
'val_rmse': 663.5009430476919,
'r_2_score': 0.9764102211643134,
'Model Accuracy': 97.64102211643134}

```

```

model1.get_importance_df()

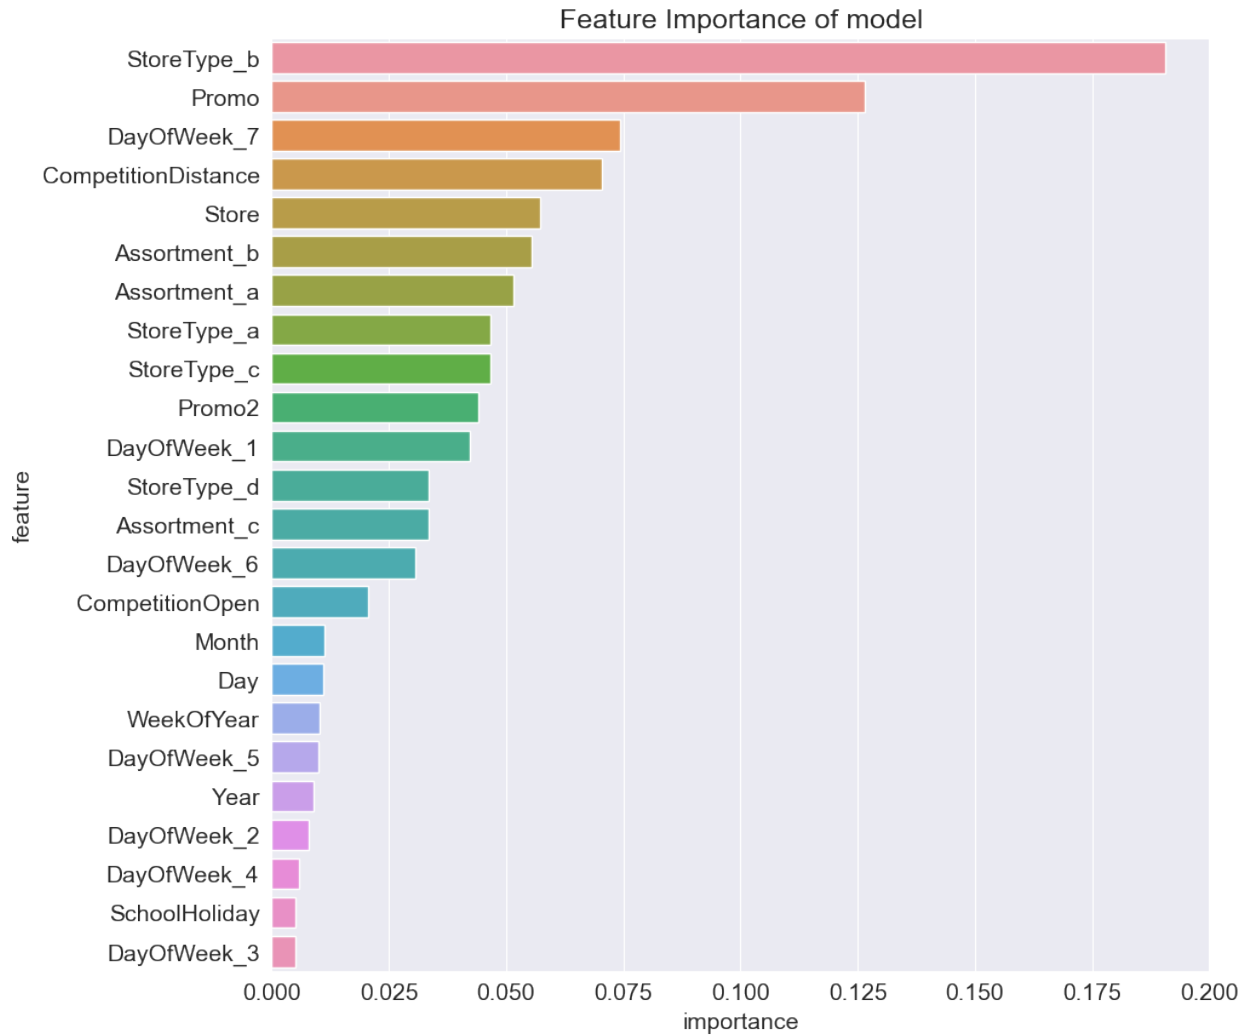
```

	feature	importance
18	StoreType_b	0.190685
1	Promo	0.126486
16	DayOfWeek_7	0.074340
3	CompetitionDistance	0.070415
0	Store	0.057211
22	Assortment_b	0.055504
21	Assortment_a	0.051471
17	StoreType_a	0.046658
19	StoreType_c	0.046658
5	Promo2	0.044138
10	DayOfWeek_1	0.042414
20	StoreType_d	0.033610
23	Assortment_c	0.033391
15	DayOfWeek_6	0.030780
4	CompetitionOpen	0.020668
7	Month	0.011345
6	Day	0.010923
9	WeekOfYear	0.010193
14	DayOfWeek_5	0.009941
8	Year	0.008853
11	DayOfWeek_2	0.008038
13	DayOfWeek_4	0.005864
2	SchoolHoliday	0.005216
12	DayOfWeek_3	0.005199

```

model1.plot_importance()

```



```
new_input = {  
}  
  
model1.list_params()  
  
['train_inputs',  
 'test_inputs',  
 'train_target',  
 'val_target',  
 'numeric_cols',  
 'categorical_cols',  
 'encoded_cols',  
 'imputer_statistics',  
 'X_train',  
 'X_val',  
 'X_test',  
 'encoder',
```

```

'scaler',
'Model',
'train_rmse',
'val_rmse',
'r_2_score',
'Model Accuracy']

model_paarms= model1.get_params()
model_paarms.get('Model')

XGBRegressor(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=0.7, device=None,
early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None,
feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=0.1,
max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=10, max_leaves=None,
              min_child_weight=None, missing=nan,
monotone_constraints=None,
              multi_strategy=None, n_estimators=1000, n_jobs=-1,
              num_parallel_tree=None, random_state=42, ...)

model1.import_model()

```

Scaling Numeric Features

Another good practice is to scale numeric features to a small range of values e.g. $(0, 1)$ or $(-1, 1)$. Scaling numeric features ensures that no particular feature has a disproportionate impact on the model's loss. Optimization algorithms also work better in practice with smaller numbers.

The numeric columns in our dataset have varying ranges.

```

from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()

scaler.fit(train_inputs[numeric_cols])

MinMaxScaler()

train_inputs[numeric_cols] =
scaler.transform(train_inputs[numeric_cols])
test_inputs[numeric_cols] =
scaler.transform(test_inputs[numeric_cols])

train_inputs[numeric_cols].describe()

```

	Store	Promo	SchoolHoliday
CompetitionDistance \			
count	844392.000000	844392.000000	844392.000000
844392.000000			
mean	0.500380	0.446352	0.193580
0.074107			
std	0.288808	0.497114	0.395103
0.113142			
min	0.000000	0.000000	0.000000
0.000000			
25%	0.250449	0.000000	0.000000
0.009098			
50%	0.500000	0.000000	0.000000
0.030459			
75%	0.750449	1.000000	0.000000
0.090849			
max	1.000000	1.000000	1.000000
1.000000			

	CompetitionOpen	Promo2	Day	Month \
count	844392.000000	844392.000000	844392.000000	844392.000000
mean	0.030270	0.498684	0.494523	0.440522
std	0.047034	0.499999	0.289449	0.302176
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.233333	0.181818
50%	0.011544	0.000000	0.500000	0.454545
75%	0.052670	1.000000	0.733333	0.636364
max	1.000000	1.000000	1.000000	1.000000

	Year	WeekOfYear
count	844392.000000	844392.000000
mean	0.415969	0.444055
std	0.388630	0.282153
min	0.000000	0.000000
25%	0.000000	0.196078
50%	0.500000	0.431373
75%	0.500000	0.666667
max	1.000000	1.000000

test_inputs[numeric_cols]

	Store	Promo	SchoolHoliday	CompetitionDistance
CompetitionOpen \				
0	0.000000	1.0	0.0	0.016482
0.060606				
1	0.001795	1.0	0.0	0.186050
0.075758				
2	0.005386	1.0	0.0	0.316192
0.020924				
3	0.006284	1.0	0.0	0.098892

```

0.007937
4      0.007181      1.0      0.0      0.026503
0.130592
...      ...      ...      ...      ...
...
41083  0.996409      0.0      0.0      0.024789
0.010101
41084  0.997307      0.0      0.0      0.024525
0.080808
41085  0.998205      0.0      0.0      0.121835
0.000000
41086  0.999102      0.0      0.0      0.011208
0.000000
41087  1.000000      0.0      1.0      0.070280
0.000000

```

	Promo2	Day	Month	Year	WeekOfYear
0	0.0	0.533333	0.727273	1.0	0.725490
1	1.0	0.533333	0.727273	1.0	0.725490
2	0.0	0.533333	0.727273	1.0	0.725490
3	0.0	0.533333	0.727273	1.0	0.725490
4	0.0	0.533333	0.727273	1.0	0.725490
...
41083	1.0	0.000000	0.636364	1.0	0.588235
41084	0.0	0.000000	0.636364	1.0	0.588235
41085	0.0	0.000000	0.636364	1.0	0.588235
41086	0.0	0.000000	0.636364	1.0	0.588235
41087	1.0	0.000000	0.636364	1.0	0.588235

```
[41088 rows x 10 columns]
```

Encoding Categorical Data

Since machine learning models can only be trained with numeric data, we need to convert categorical data to numbers. A common technique is to use one-hot encoding for categorical columns.

One hot encoding involves adding a new binary (0/1) column for each unique category of a categorical column.

```

from sklearn.preprocessing import OneHotEncoder
train_inputs['StateHoliday'].unique()
array(['0', 'a', 'b', 'c', 0], dtype=object)
train_inputs[categorical_cols]

```

	DayOfWeek	StoreType	Assortment
0	5	c	a
1	5	a	a
2	5	a	a
3	5	c	c
4	5	a	a
...
1016776	2	b	a
1016827	2	b	b
1016863	2	b	b
1017042	2	b	b
1017190	2	b	b

[844392 rows x 3 columns]

```
encoder = OneHotEncoder(sparse=False,
handle_unknown='ignore').fit(train_inputs[categorical_cols])
encoded_cols = list(encoder.get_feature_names_out(categorical_cols))
```

```
train_inputs[encoded_cols] =
encoder.transform(train_inputs[categorical_cols])
test_inputs[encoded_cols] =
encoder.transform(test_inputs[categorical_cols])
```

```
X_train = train_inputs[numeric_cols + encoded_cols]
X_test = test_inputs[numeric_cols + encoded_cols]
```

X_train

	Store	Promo	SchoolHoliday	CompetitionDistance
CompetitionOpen \				
0	0.000000	1.0	1.0	0.016482
0.059163				
1	0.000898	1.0	1.0	0.007252
0.066378				
2	0.001795	1.0	1.0	0.186050
0.074315				
3	0.002693	1.0	1.0	0.007911
0.050505				
4	0.003591	1.0	1.0	0.394119
0.002165				
...
...				
1016776	0.611311	0.0	1.0	0.001714
0.054834				
1016827	0.657092	0.0	1.0	0.011076
0.114719				
1016863	0.689408	0.0	1.0	0.010812
0.000000				
1017042	0.850090	0.0	1.0	0.018592

0.000000

1017190 0.983842 0.0 1.0 0.009230

0.093795

	Promo2	Day	Month	Year	WeekOfYear	...	DayOfWeek_5 \
0	0.0	1.0	0.545455	1.0	0.588235	...	1.0
1	1.0	1.0	0.545455	1.0	0.588235	...	1.0
2	1.0	1.0	0.545455	1.0	0.588235	...	1.0
3	0.0	1.0	0.545455	1.0	0.588235	...	1.0
4	0.0	1.0	0.545455	1.0	0.588235	...	1.0
...
1016776	0.0	0.0	0.000000	0.0	0.000000	...	0.0
1016827	0.0	0.0	0.000000	0.0	0.000000	...	0.0
1016863	1.0	0.0	0.000000	0.0	0.000000	...	0.0
1017042	0.0	0.0	0.000000	0.0	0.000000	...	0.0
1017190	0.0	0.0	0.000000	0.0	0.000000	...	0.0

	DayOfWeek_6	DayOfWeek_7	StoreType_a	StoreType_b
StoreType_c \				
0	0.0	0.0	0.0	0.0
1.0				
1	0.0	0.0	1.0	0.0
0.0				
2	0.0	0.0	1.0	0.0
0.0				
3	0.0	0.0	0.0	0.0
1.0				
4	0.0	0.0	1.0	0.0
0.0				
...
...				
1016776	0.0	0.0	0.0	1.0
0.0				
1016827	0.0	0.0	0.0	1.0
0.0				
1016863	0.0	0.0	0.0	1.0
0.0				
1017042	0.0	0.0	0.0	1.0
0.0				
1017190	0.0	0.0	0.0	1.0
0.0				

	StoreType_d	Assortment_a	Assortment_b	Assortment_c
0	0.0	1.0	0.0	0.0
1	0.0	1.0	0.0	0.0
2	0.0	1.0	0.0	0.0
3	0.0	0.0	0.0	1.0
4	0.0	1.0	0.0	0.0
...
1016776	0.0	1.0	0.0	0.0

1016827	0.0	0.0	1.0	0.0
1016863	0.0	0.0	1.0	0.0
1017042	0.0	0.0	1.0	0.0
1017190	0.0	0.0	1.0	0.0

[844392 rows x 24 columns]

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_val, train_targets, val_targets = train_test_split(X_train,
train_targets, test_size=0.1)
```

X_train

	Store	Promo	SchoolHoliday	CompetitionDistance
CompetitionOpen \				
461690	0.776481	0.0	0.0	0.127373
0.000000				
58084	0.093357	0.0	0.0	0.081355
0.000000				
936213	0.357271	0.0	0.0	0.070411
0.004329				
535814	0.254937	0.0	0.0	0.031514
0.000000				
386157	0.197487	0.0	1.0	0.178138
0.007215				
...
...				
210968	0.209156	0.0	0.0	0.057358
0.000000				
418832	0.338420	0.0	0.0	0.027954
0.015873				
538095	0.300718	1.0	0.0	0.002242
0.000000				
791195	0.296230	0.0	1.0	0.008571
0.000000				
768538	0.976661	1.0	1.0	0.068565
0.036797				

	Promo2	Day	Month	Year	WeekOfYear	...	DayOfWeek_5
\							
461690	1.0	0.433333	0.363636	0.5	0.372549	...	0.0
58084	1.0	0.266667	0.454545	1.0	0.450980	...	0.0
936213	1.0	0.433333	0.181818	0.0	0.196078	...	0.0
535814	0.0	0.233333	0.181818	0.5	0.176471	...	0.0
386157	0.0	0.766667	0.545455	0.5	0.568627	...	0.0

...
210968	0.0	0.733333	0.000000	1.0	0.058824	...	1.0
418832	0.0	0.666667	0.454545	0.5	0.470588	...	0.0
538095	0.0	0.166667	0.181818	0.5	0.176471	...	0.0
791195	1.0	0.700000	0.545455	0.0	0.568627	...	0.0
768538	0.0	0.366667	0.636364	0.0	0.627451	...	0.0
	DayOfWeek_6	DayOfWeek_7	StoreType_a	StoreType_b			
StoreType_c \							
461690	0.0	0.0	0.0	0.0			
0.0							
58084	0.0	0.0	1.0	0.0			
0.0							
936213	0.0	0.0	1.0	0.0			
0.0							
535814	1.0	0.0	1.0	0.0			
0.0							
386157	0.0	0.0	0.0	0.0			
0.0							
...
.							
210968	0.0	0.0	0.0	0.0			
0.0							
418832	1.0	0.0	1.0	0.0			
0.0							
538095	0.0	0.0	1.0	0.0			
0.0							
791195	0.0	0.0	1.0	0.0			
0.0							
768538	0.0	0.0	0.0	0.0			
0.0							
	StoreType_d	Assortment_a	Assortment_b	Assortment_c			
461690	1.0	1.0	0.0	0.0			
58084	0.0	0.0	0.0	1.0			
936213	0.0	1.0	0.0	0.0			
535814	0.0	1.0	0.0	0.0			
386157	1.0	0.0	0.0	1.0			
...			
210968	1.0	1.0	0.0	0.0			
418832	0.0	0.0	0.0	1.0			
538095	0.0	1.0	0.0	0.0			
791195	0.0	0.0	0.0	1.0			
768538	1.0	1.0	0.0	0.0			

[759952 rows x 24 columns]

X_val

	Store	Promo	SchoolHoliday	CompetitionDistance
CompetitionOpen \				
167186	0.943447	1.0	0.0	0.066719
0.000000				
896616	0.844704	0.0	0.0	0.090190
0.000000				
460211	0.449731	0.0	0.0	0.002637
0.106061				
430237	0.567325	0.0	1.0	0.153217
0.079365				
552682	0.383303	1.0	0.0	0.038766
0.000000				
...
...				
2628	0.357271	1.0	0.0	0.070411
0.024531				
952982	0.396768	0.0	0.0	0.150053
0.062049				
973111	0.449731	0.0	0.0	0.002637
0.095238				
661273	0.774686	0.0	0.0	0.013186
0.009380				
327890	0.874327	0.0	0.0	0.126714
0.000000				

	Promo2	Day	Month	Year	WeekOfYear	...	DayOfWeek_5
\							
167186	1.0	0.100000	0.181818	1.0	0.176471	...	0.0
896616	1.0	0.600000	0.272727	0.0	0.294118	...	1.0
460211	1.0	0.466667	0.363636	0.5	0.372549	...	0.0
430237	1.0	0.333333	0.454545	0.5	0.450980	...	0.0
552682	1.0	0.666667	0.090909	0.5	0.137255	...	1.0
...
2628	1.0	0.933333	0.545455	1.0	0.588235	...	0.0
952982	0.0	0.866667	0.090909	0.0	0.156863	...	0.0
973111	1.0	0.266667	0.090909	0.0	0.098039	...	0.0
661273	1.0	0.500000	0.909091	0.0	0.882353	...	0.0

327890	1.0	0.800000	0.727273	0.5	0.745098	...	0.0
--------	-----	----------	----------	-----	----------	-----	-----

	DayOfWeek_6	DayOfWeek_7	StoreType_a	StoreType_b
StoreType_c \				
167186	0.0	0.0	1.0	0.0
0.0				
896616	0.0	0.0	0.0	0.0
0.0				
460211	0.0	0.0	1.0	0.0
0.0				
430237	0.0	0.0	0.0	0.0
0.0				
552682	0.0	0.0	0.0	0.0
0.0				
...
.				
2628	0.0	0.0	1.0	0.0
0.0				
952982	0.0	0.0	0.0	0.0
0.0				
973111	1.0	0.0	1.0	0.0
0.0				
661273	1.0	0.0	1.0	0.0
0.0				
327890	0.0	0.0	1.0	0.0
0.0				

	StoreType_d	Assortment_a	Assortment_b	Assortment_c
167186	0.0	0.0	0.0	1.0
896616	1.0	0.0	0.0	1.0
460211	0.0	1.0	0.0	0.0
430237	1.0	1.0	0.0	0.0
552682	1.0	1.0	0.0	0.0
...
2628	0.0	1.0	0.0	0.0
952982	1.0	1.0	0.0	0.0
973111	0.0	1.0	0.0	0.0
661273	0.0	1.0	0.0	0.0
327890	0.0	0.0	0.0	1.0

[84440 rows x 24 columns]

```
def rmse(targets , predictions):
    return np.sqrt(np.mean(np.square(targets-predictions)))

def mse(targets , predictions):
    return np.mean(np.square(targets-predictions))
```

Gradient Boosting

We're now ready to train our gradient boosting machine (GBM) model. Here's how a GBM model works:

1. The average value of the target column and uses as an initial prediction every input.
2. The residuals (difference) of the predictions with the targets are computed.
3. A decision tree of limited depth is trained to **predict just the residuals** for each input.
4. Predictions from the decision tree are scaled using a parameter called the learning rate (this prevents overfitting)
5. Scaled predictions for the tree are added to the previous predictions to obtain the new and improved predictions.
6. Steps 2 to 5 are repeated to create new decision trees, each of which is trained to predict just the residuals from the previous prediction.

The term "gradient" refers to the fact that each decision tree is trained with the purpose of reducing the loss from the previous iteration (similar to gradient descent). The term "boosting" refers the general technique of training new models to improve the results of an existing model.

EXERCISE: Can you describe in your own words how a gradient boosting machine is different from a random forest?

For a mathematical explanation of gradient boosting, check out the following resources:

- [XGBoost Documentation](#)
- [Video Tutorials on StatQuest](#)

Here's a visual representation of gradient boosting:

```
from sklearn.metrics import r2_score ,accuracy_score
from xgboost import XGBRegressor
```

```

## Model Training and Hyper parameter Tuning
xgb_model = XGBRegressor(n_jobs=-1, random_state=42,
n_estimators=1000,
                        learning_rate=0.1, max_depth=10, subsample=0.9,
                        colsample_bytree=0.7)
# Fiting Model to Dataset -- X_train, train_targets
xgb_model.fit(X_train,train_targets)

# Making Prediction
train_preds = xgb_model.predict(X_train)
val_preds = xgb_model.predict(X_val)
r_2_score = xgb_model.score(X_train,train_targets)
accuracy = xgb_model.score(X_train,train_targets)*100

print(xgb_model)
print('Training RMSE : ',rmse(train_targets,train_preds))
print('Testing RMSE : ',rmse(val_targets,val_preds))
print('Model R_2 score : ',r_2_score)
print('Model Accuracy : ',accuracy)

XGBRegressor(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=0.7, device=None,
early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None,
feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=0.1,
max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=10, max_leaves=None,
              min_child_weight=None, missing=nan,
monotone_constraints=None,
              multi_strategy=None, n_estimators=1000, n_jobs=-1,
              num_parallel_tree=None, random_state=42, ...)
Training RMSE : 478.61037819807126
Testing RMSE : 671.8424331144322
Model R_2 score : 0.9762329659464672
Model Accuracy : 97.62329659464672

importance_xgbregressor_df = pd.DataFrame({
    'feature': X_train.columns,
    'importance': xgb_model.feature_importances_
}).sort_values('importance', ascending=False)

importance_xgbregressor_df.head()


```

	feature	importance
18	StoreType_b	0.184327
1	Promo	0.125972

```
16         DayOfWeek_7      0.077401
3  CompetitionDistance      0.067803
22        Assortment_b      0.060906
```

```
from sklearn.metrics import r2_score
```

```
plt.figure(figsize=(10,10))
plt.title('Feature Importance xgBoost Regressor')
sns.barplot(data=importance_xgbregressor_df, x='importance',
y='feature')
plt.show()
```



Decision Tree Regressor

Decision Trees (DTs) are a non-parametric supervised learning method used for classification and regression. The goal is to create a model that predicts the value of a target variable by

learning simple decision rules inferred from the data features. A tree can be seen as a piecewise constant approximation.

For instance, in the example below, decision trees learn from data to approximate a sine curve with a set of if-then-else decision rules. The deeper the tree, the more complex the decision rules and the fitter the model.

Some advantages of decision trees are:

Simple to understand and to interpret. Trees can be visualized.

Requires little data preparation. Other techniques often require data normalization, dummy variables need to be created and blank values to be removed. Some tree and algorithm combinations support missing values.

The cost of using the tree (i.e., predicting data) is logarithmic in the number of data points used to train the tree.

Able to handle both numerical and categorical data. However, the scikit-learn implementation does not support categorical variables for now. Other techniques are usually specialized in analyzing datasets that have only one type of variable. See algorithms for more information.

Able to handle multi-output problems.

Uses a white box model. If a given situation is observable in a model, the explanation for the condition is easily explained by boolean logic. By contrast, in a black box model (e.g., in an artificial neural network), results may be more difficult to interpret.

Possible to validate a model using statistical tests. That makes it possible to account for the reliability of the model.

Performs well even if its assumptions are somewhat violated by the true model from which the data were generated.

The disadvantages of decision trees include:

Decision-tree learners can create over-complex trees that do not generalize the data well. This is called overfitting. Mechanisms such as pruning, setting the minimum number of samples required at a leaf node or setting the maximum depth of the tree are necessary to avoid this problem.

Decision trees can be unstable because small variations in the data might result in a completely different tree being generated. This problem is mitigated by using decision trees within an ensemble.

Predictions of decision trees are neither smooth nor continuous, but piecewise constant approximations as seen in the above figure. Therefore, they are not good at extrapolation.

The problem of learning an optimal decision tree is known to be NP-complete under several aspects of optimality and even for simple concepts. Consequently, practical decision-tree learning algorithms are based on heuristic algorithms such as the greedy algorithm where locally optimal decisions are made at each node. Such algorithms cannot guarantee to return the

globally optimal decision tree. This can be mitigated by training multiple trees in an ensemble learner, where the features and samples are randomly sampled with replacement.

There are concepts that are hard to learn because decision trees do not express them easily, such as XOR, parity or multiplexer problems.

Decision tree learners create biased trees if some classes dominate. It is therefore recommended to balance the dataset prior to fitting with the decision tree

```
from sklearn.tree import DecisionTreeRegressor
```

?DecisionTreeRegressor

Init signature:

```
DecisionTreeRegressor(  
    *,  
    criterion='squared_error',  
    splitter='best',  
    max_depth=None,  
    min_samples_split=2,  
    min_samples_leaf=1,  
    min_weight_fraction_leaf=0.0,  
    max_features=None,  
    random_state=None,  
    max_leaf_nodes=None,  
    min_impurity_decrease=0.0,  
    ccp_alpha=0.0,  
)
```

Docstring:

A decision tree regressor.

Read more in the :ref:`User Guide <tree>`.

Parameters

criterion : {"squared_error", "friedman_mse", "absolute_error", "poisson"}, default="squared_error"

The function to measure the quality of a split. Supported criteria are "squared_error" for the mean squared error, which is equal to variance reduction as feature selection criterion and minimizes the L2

loss using the mean of each terminal node, "friedman_mse", which uses

mean squared error with Friedman's improvement score for potential splits, "absolute_error" for the mean absolute error, which minimizes

the L1 loss using the median of each terminal node, and "poisson" which

uses reduction in Poisson deviance to find splits.

.. versionadded:: 0.18

Mean Absolute Error (MAE) criterion.

.. versionadded:: 0.24

Poisson deviance criterion.

`splitter : {"best", "random"}, default="best"`

The strategy used to choose the split at each node. Supported strategies are "best" to choose the best split and "random" to choose the best random split.

`max_depth : int, default=None`

The maximum depth of the tree. If None, then nodes are expanded until all leaves are pure or until all leaves contain less than `min_samples_split` samples.

`min_samples_split : int or float, default=2`

The minimum number of samples required to split an internal node:

- If int, then consider ``min_samples_split`` as the minimum number.
- If float, then ``min_samples_split`` is a fraction and ``ceil(min_samples_split * n_samples)`` are the minimum number of samples for each split.

.. versionchanged:: 0.18

Added float values for fractions.

`min_samples_leaf : int or float, default=1`

The minimum number of samples required to be at a leaf node. A split point at any depth will only be considered if it leaves at least ``min_samples_leaf`` training samples in each of the left and right branches. This may have the effect of smoothing the model, especially in regression.

- If int, then consider ``min_samples_leaf`` as the minimum number.
- If float, then ``min_samples_leaf`` is a fraction and ``ceil(min_samples_leaf * n_samples)`` are the minimum number of samples for each node.

.. versionchanged:: 0.18

Added float values for fractions.

`min_weight_fraction_leaf : float, default=0.0`

The minimum weighted fraction of the sum total of weights (of all the input samples) required to be at a leaf node. Samples have equal weight when `sample_weight` is not provided.

`max_features : int, float or {"auto", "sqrt", "log2"}, default=None`

The number of features to consider when looking for the best split:

- If int, then consider ``max_features`` features at each split.
- If float, then ``max_features`` is a fraction and ``max(1, int(max_features * n_features_in_))`` features are considered at each split.
- If "sqrt", then ``max_features=sqrt(n_features)``.
- If "log2", then ``max_features=log2(n_features)``.
- If None, then ``max_features=n_features``.

Note: the search for a split does not stop until at least one valid partition of the node samples is found, even if it requires to effectively inspect more than ``max_features`` features.

`random_state` : int, RandomState instance or None, default=None
Controls the randomness of the estimator. The features are always randomly permuted at each split, even if ``splitter`` is set to ``"best"``. When ``max_features < n_features``, the algorithm will select ``max_features`` at random at each split before finding the best split among them. But the best found split may vary across different runs, even if ``max_features=n_features``. That is the case, if the improvement of the criterion is identical for several splits and one split has to be selected at random. To obtain a deterministic behaviour during fitting, ``random_state`` has to be fixed to an integer. See :term:`Glossary <random_state>` for details.

`max_leaf_nodes` : int, default=None
Grow a tree with ``max_leaf_nodes`` in best-first fashion. Best nodes are defined as relative reduction in impurity. If None then unlimited number of leaf nodes.

`min_impurity_decrease` : float, default=0.0
A node will be split if this split induces a decrease of the impurity greater than or equal to this value.

The weighted impurity decrease equation is the following::

$$N_t / N * (impurity - N_{t_R} / N_t * right_impurity - N_{t_L} / N_t * left_impurity)$$

where ``N`` is the total number of samples, ``N_t`` is the number

of
samples at the current node, ``N_t_L`` is the number of samples in
the
left child, and ``N_t_R`` is the number of samples in the right
child.

``N``, ``N_t``, ``N_t_R`` and ``N_t_L`` all refer to the weighted
sum,
if ``sample_weight`` is passed.

.. versionadded:: 0.19

ccp_alpha : non-negative float, default=0.0
Complexity parameter used for Minimal Cost-Complexity Pruning. The
subtree with the largest cost complexity that is smaller than
``ccp_alpha`` will be chosen. By default, no pruning is performed.

See
:ref:`minimal_cost_complexity_pruning` for details.

.. versionadded:: 0.22

Attributes

feature_importances_ : ndarray of shape (n_features,)

The feature importances.

The higher, the more important the feature.

The importance of a feature is computed as the
(normalized) total reduction of the criterion brought
by that feature. It is also known as the Gini importance [4]_.

Warning: impurity-based feature importances can be misleading for
high cardinality features (many unique values). See
:func:`sklearn.inspection.permutation_importance` as an
alternative.

max_features_ : int

The inferred value of max_features.

n_features_in_ : int

Number of features seen during :term:`fit`.

.. versionadded:: 0.24

feature_names_in_ : ndarray of shape (n_features_in_,)

Names of features seen during :term:`fit`. Defined only when `X`
has feature names that are all strings.

.. versionadded:: 1.0

n_outputs_ : int

The number of outputs when ``fit`` is performed.

`tree_` : Tree instance

The underlying Tree object. Please refer to

``help(sklearn.tree._tree.Tree)`` for attributes of Tree object
and

:ref:`sphx_glr_auto_examples_tree_plot_unveil_tree_structure.py`
for basic usage of these attributes.

See Also

`DecisionTreeClassifier` : A decision tree classifier.

Notes

The default values for the parameters controlling the size of the
trees

(e.g. ``max_depth``, ``min_samples_leaf``, etc.) lead to fully grown
and

unpruned trees which can potentially be very large on some data sets.

To

reduce memory consumption, the complexity and size of the trees should
be

controlled by setting those parameter values.

References

.. [1] https://en.wikipedia.org/wiki/Decision_tree_learning

.. [2] L. Breiman, J. Friedman, R. Olshen, and C. Stone,
"Classification
and Regression Trees", Wadsworth, Belmont, CA, 1984.

.. [3] T. Hastie, R. Tibshirani and J. Friedman. "Elements of
Statistical
Learning", Springer, 2009.

.. [4] L. Breiman, and A. Cutler, "Random Forests",
https://www.stat.berkeley.edu/~breiman/RandomForests/cc_home.htm

Examples

```
>>> from sklearn.datasets import load_diabetes
>>> from sklearn.model_selection import cross_val_score
>>> from sklearn.tree import DecisionTreeRegressor
>>> X, y = load_diabetes(return_X_y=True)
>>> regressor = DecisionTreeRegressor(random_state=0)
>>> cross_val_score(regressor, X, y, cv=10)
```

```

...                                     # doctest: +SKIP
...
array([-0.39..., -0.46...,  0.02...,  0.06..., -0.50...,
       0.16...,  0.11..., -0.73..., -0.30..., -0.00...])
File:      c:\users\saket\appdata\local\programs\python\
python311\lib\site-packages\sklearn\tree\_classes.py
Type:      ABCMeta
Subclasses: ExtraTreeRegressor

## Model Training and Hyper parameter Tuning
decision_tree_model = DecisionTreeRegressor(random_state=12 ,
max_depth=12 , max_leaf_nodes=2**20 ,min_samples_split=15 )
# Fiting Model to Dataset -- X_train, train_targets
decision_tree_model.fit(X_train,train_targets)

# Making Prediction
train_preds = decision_tree_model.predict(X_train)
val_preds = decision_tree_model.predict(X_val)
r_2_score = decision_tree_model.score(X_train,train_targets)
accuracy = decision_tree_model.score(X_train,train_targets)*100

print(decision_tree_model)
print('Training RMSE : ',rmse(train_targets,train_preds))
print('Testing RMSE : ',rmse(val_targets,val_preds))
print('Model r_2 score : ',r_2_score)
print('Model Accuracy : ',accuracy)

DecisionTreeRegressor(max_depth=12, max_leaf_nodes=1048576,
                      min_samples_split=15, random_state=12)
Training RMSE :  2132.7798691120297
Testing RMSE :  2173.2814436015874
Model r_2 score :  0.5280425663493122
Model Accuracy :  52.80425663493122

decision_tree_model.feature_importances_

array([1.43192757e-01, 2.56351287e-01, 6.54033381e-03, 2.41440344e-01,
       3.27276444e-02, 4.66660024e-02, 3.02306257e-02, 8.82616074e-03,
       3.26795741e-03, 3.92584852e-02, 5.63993511e-02, 2.08982019e-03,
       4.11403487e-05, 1.47682764e-04, 3.47045767e-03, 9.28093906e-03,
       7.95077941e-03, 7.90502633e-03, 4.59897490e-02, 1.29668663e-02,
       5.79147029e-03, 1.09401279e-02, 2.06858053e-03, 2.64564109e-
02])


importance_decision_tree_df = pd.DataFrame({
    'feature': X_train.columns,
    'importance': decision_tree_model.feature_importances_
}).sort_values('importance', ascending=False)

importance_decision_tree_df.head(5)

```

	feature	importance
1	Promo	0.256351
3	CompetitionDistance	0.241440
0	Store	0.143193
10	DayOfWeek_1	0.056399
5	Promo2	0.046666

```
plt.figure(figsize=(10,10))
plt.title('Feature Importance Decision Tree Regressor')
sns.barplot(data=importance_decision_tree_df, x='importance',
y='feature')
plt.show()
```



Saving and Loading Trained Models

We can save the parameters (weights and biases) of our trained model to disk, so that we needn't retrain the model from scratch each time we wish to use it. Along with the model, it's

also important to save imputers, scalers, encoders and even column names. Anything that will be required while generating predictions using the model should be saved.

We can use the `joblib` module to save and load Python objects on the disk.

```
import joblib

rossman_xgboost_model = {
    'model': xgb_model,
    'scaler': scaler,
    'encoder': encoder,
    'input_cols': input_cols,
    'train_inputs' : train_inputs,
    'train_targets' : train_targets,
    'test_inputs' : test_inputs,
    'X_train' : X_train,
    'X_test' : X_test,
    'target_col': target_col,
    'numeric_cols': numeric_cols,
    'categorical_cols': categorical_cols,
    'encoded_cols': encoded_cols
}

joblib.dump(rossman_xgboost_model, "files\\
rossman_xgboost_model.joblib")

['files\\rossman_xgboost_model.joblib']

joblib.load("files\\rossman_xgboost_model.joblib")

{'model': XGBRegressor(base_score=None, booster=None, callbacks=None,
    colsample_bylevel=None, colsample_bynode=None,
    colsample_bytree=0.7, device=None,
early_stopping_rounds=None,
    enable_categorical=False, eval_metric=None,
feature_types=None,
    gamma=None, grow_policy=None, importance_type=None,
    interaction_constraints=None, learning_rate=0.1,
max_bin=None,
    max_cat_threshold=None, max_cat_to_onehot=None,
    max_delta_step=None, max_depth=10, max_leaves=None,
    min_child_weight=None, missing=nan,
monotone_constraints=None,
    multi_strategy=None, n_estimators=1000, n_jobs=-1,
    num_parallel_tree=None, random_state=42, ...),
'scaler': MinMaxScaler(),
'encoder': OneHotEncoder(handle_unknown='ignore', sparse=False,
sparse_output=False),
'input_cols': ['Store',
'DayOfWeek',
'Promo',
```

```

'StateHoliday',
'Assortment',
'CompetitionDistance',
'Day',
'Month',
'Year',
'WeekOfYear',
'Promo2',
'CompetitionOpen',
'SchoolHoliday',
'StoreType'],
'train_inputs':

```

		Store	DayOfWeek	Promo	StateHoliday
Assortment \					
0	0.000000	5	1.0	0	a
1	0.000898	5	1.0	0	a
2	0.001795	5	1.0	0	a
3	0.002693	5	1.0	0	c
4	0.003591	5	1.0	0	a
...
1016776	0.611311	2	0.0	a	a
1016827	0.657092	2	0.0	a	b
1016863	0.689408	2	0.0	a	b
1017042	0.850090	2	0.0	a	b
1017190	0.983842	2	0.0	a	b

	CompetitionDistance	Day	Month	Year	WeekOfYear	...	\
0	0.016482	1.0	0.545455	1.0	0.588235	...	
1	0.007252	1.0	0.545455	1.0	0.588235	...	
2	0.186050	1.0	0.545455	1.0	0.588235	...	
3	0.007911	1.0	0.545455	1.0	0.588235	...	
4	0.394119	1.0	0.545455	1.0	0.588235	...	
...	
1016776	0.001714	0.0	0.000000	0.0	0.000000	...	
1016827	0.011076	0.0	0.000000	0.0	0.000000	...	
1016863	0.010812	0.0	0.000000	0.0	0.000000	...	
1017042	0.018592	0.0	0.000000	0.0	0.000000	...	
1017190	0.009230	0.0	0.000000	0.0	0.000000	...	

	DayOfWeek_5	DayOfWeek_6	DayOfWeek_7	StoreType_a
StoreType_b \				
0	1.0	0.0	0.0	0.0
0.0				
1	1.0	0.0	0.0	1.0
0.0				
2	1.0	0.0	0.0	1.0
0.0				
3	1.0	0.0	0.0	0.0
0.0				
4	1.0	0.0	0.0	1.0


```

0.0
...      ...      ...      ...      ...
..
1016776      0.0      0.0      0.0      0.0
1.0
1016827      0.0      0.0      0.0      0.0
1.0
1016863      0.0      0.0      0.0      0.0
1.0
1017042      0.0      0.0      0.0      0.0
1.0
1017190      0.0      0.0      0.0      0.0
1.0

      StoreType_c  StoreType_d  Assortment_a  Assortment_b
Assortment_c
0      1.0      0.0      1.0      0.0
0.0
1      0.0      0.0      1.0      0.0
0.0
2      0.0      0.0      1.0      0.0
0.0
3      1.0      0.0      0.0      0.0
1.0
4      0.0      0.0      1.0      0.0
0.0
...      ...      ...      ...      ...
...
1016776      0.0      0.0      1.0      0.0
0.0
1016827      0.0      0.0      0.0      1.0
0.0
1016863      0.0      0.0      0.0      1.0
0.0
1017042      0.0      0.0      0.0      1.0
0.0
1017190      0.0      0.0      0.0      1.0
0.0

[844392 rows x 28 columns],
'train_targets': 461690      4360
58084      5154
936213      4369
535814      5943
386157      4837
...
210968      8291
418832      6057
538095      15623

```

```

791195      5429
768538      13436
Name: Sales, Length: 759952, dtype: int64,
'test_inputs':      Store DayOfWeek Promo StateHoliday
Assortment \
0      0.000000      4      1.0      0      a
1      0.001795      4      1.0      0      a
2      0.005386      4      1.0      0      c
3      0.006284      4      1.0      0      a
4      0.007181      4      1.0      0      c
...      ...      ...      ...      ...      ...
41083  0.996409      6      0.0      0      a
41084  0.997307      6      0.0      0      c
41085  0.998205      6      0.0      0      c
41086  0.999102      6      0.0      0      c
41087  1.000000      6      0.0      0      c

      CompetitionDistance      Day      Month      Year      WeekOfYear      ...
\
0      0.016482  0.533333  0.727273      1.0      0.725490      ...
1      0.186050  0.533333  0.727273      1.0      0.725490      ...
2      0.316192  0.533333  0.727273      1.0      0.725490      ...
3      0.098892  0.533333  0.727273      1.0      0.725490      ...
4      0.026503  0.533333  0.727273      1.0      0.725490      ...
...      ...      ...      ...      ...      ...      ...
41083      0.024789  0.000000  0.636364      1.0      0.588235      ...
41084      0.024525  0.000000  0.636364      1.0      0.588235      ...
41085      0.121835  0.000000  0.636364      1.0      0.588235      ...
41086      0.011208  0.000000  0.636364      1.0      0.588235      ...
41087      0.070280  0.000000  0.636364      1.0      0.588235      ...

      DayOfWeek_5      DayOfWeek_6      DayOfWeek_7      StoreType_a      StoreType_b
\
0      0.0      0.0      0.0      0.0      0.0
1      0.0      0.0      0.0      1.0      0.0
2      0.0      0.0      0.0      1.0      0.0
3      0.0      0.0      0.0      1.0      0.0

```

4	0.0	0.0	0.0	1.0	0.0
...
41083	0.0	1.0	0.0	1.0	0.0
41084	0.0	1.0	0.0	0.0	0.0
41085	0.0	1.0	0.0	1.0	0.0
41086	0.0	1.0	0.0	1.0	0.0
41087	0.0	1.0	0.0	0.0	0.0

	StoreType_c	StoreType_d	Assortment_a	Assortment_b
Assortment_c				
0	1.0	0.0	1.0	0.0
0.0				
1	0.0	0.0	1.0	0.0
0.0				
2	0.0	0.0	0.0	0.0
1.0				
3	0.0	0.0	1.0	0.0
0.0				
4	0.0	0.0	0.0	0.0
1.0				
...
...				
41083	0.0	0.0	1.0	0.0
0.0				
41084	1.0	0.0	0.0	0.0
1.0				
41085	0.0	0.0	0.0	0.0
1.0				
41086	0.0	0.0	0.0	0.0
1.0				
41087	0.0	1.0	0.0	0.0
1.0				

```
[41088 rows x 28 columns],
'X_train':      Store  Promo  SchoolHoliday
CompetitionDistance  CompetitionOpen  \
461690  0.776481    0.0          0.0          0.127373
0.000000
58084  0.093357    0.0          0.0          0.081355
0.000000
936213  0.357271    0.0          0.0          0.070411
0.004329
```

535814	0.254937	0.0	0.0	0.031514
0.000000				
386157	0.197487	0.0	1.0	0.178138
0.007215				
...
...				
210968	0.209156	0.0	0.0	0.057358
0.000000				
418832	0.338420	0.0	0.0	0.027954
0.015873				
538095	0.300718	1.0	0.0	0.002242
0.000000				
791195	0.296230	0.0	1.0	0.008571
0.000000				
768538	0.976661	1.0	1.0	0.068565
0.036797				

	Promo2	Day	Month	Year	WeekOfYear	...
DayOfWeek_5 \						
461690	1.0	0.433333	0.363636	0.5	0.372549	...
0.0						
58084	1.0	0.266667	0.454545	1.0	0.450980	...
0.0						
936213	1.0	0.433333	0.181818	0.0	0.196078	...
0.0						
535814	0.0	0.233333	0.181818	0.5	0.176471	...
0.0						
386157	0.0	0.766667	0.545455	0.5	0.568627	...
0.0						
...
.						
210968	0.0	0.733333	0.000000	1.0	0.058824	...
1.0						
418832	0.0	0.666667	0.454545	0.5	0.470588	...
0.0						
538095	0.0	0.166667	0.181818	0.5	0.176471	...
0.0						
791195	1.0	0.700000	0.545455	0.0	0.568627	...
0.0						
768538	0.0	0.366667	0.636364	0.0	0.627451	...
0.0						

	DayOfWeek_6	DayOfWeek_7	StoreType_a	StoreType_b
StoreType_c \				
461690	0.0	0.0	0.0	0.0
0.0				
58084	0.0	0.0	1.0	0.0
0.0				
936213	0.0	0.0	1.0	0.0

```

0.0
535814      1.0      0.0      1.0      0.0
0.0
386157      0.0      0.0      0.0      0.0
0.0
...      ...      ...      ...      .
..
210968      0.0      0.0      0.0      0.0
0.0
418832      1.0      0.0      1.0      0.0
0.0
538095      0.0      0.0      1.0      0.0
0.0
791195      0.0      0.0      1.0      0.0
0.0
768538      0.0      0.0      0.0      0.0
0.0

StoreType_d Assortment_a Assortment_b Assortment_c
461690      1.0      1.0      0.0      0.0
58084      0.0      0.0      0.0      1.0
936213      0.0      1.0      0.0      0.0
535814      0.0      1.0      0.0      0.0
386157      1.0      0.0      0.0      1.0
...      ...      ...      ...      ...
210968      1.0      1.0      0.0      0.0
418832      0.0      0.0      0.0      1.0
538095      0.0      1.0      0.0      0.0
791195      0.0      0.0      0.0      1.0
768538      1.0      1.0      0.0      0.0

[759952 rows x 24 columns],
'X_test':      Store Promo SchoolHoliday CompetitionDistance
CompetitionOpen \
0      0.000000      1.0      0.0      0.016482
0.060606
1      0.001795      1.0      0.0      0.186050
0.075758
2      0.005386      1.0      0.0      0.316192
0.020924
3      0.006284      1.0      0.0      0.098892
0.007937
4      0.007181      1.0      0.0      0.026503
0.130592
...      ...      ...      ...      ...
...
41083      0.996409      0.0      0.0      0.024789
0.010101
41084      0.997307      0.0      0.0      0.024525

```

0.080808

41085 0.998205 0.0 0.0 0.121835

0.000000

41086 0.999102 0.0 0.0 0.011208

0.000000

41087 1.000000 0.0 1.0 0.070280

0.000000

	Promo2	Day	Month	Year	WeekOfYear	...	DayOfWeek_5
\							
0	0.0	0.533333	0.727273	1.0	0.725490	...	0.0
1	1.0	0.533333	0.727273	1.0	0.725490	...	0.0
2	0.0	0.533333	0.727273	1.0	0.725490	...	0.0
3	0.0	0.533333	0.727273	1.0	0.725490	...	0.0
4	0.0	0.533333	0.727273	1.0	0.725490	...	0.0
...
41083	1.0	0.000000	0.636364	1.0	0.588235	...	0.0
41084	0.0	0.000000	0.636364	1.0	0.588235	...	0.0
41085	0.0	0.000000	0.636364	1.0	0.588235	...	0.0
41086	0.0	0.000000	0.636364	1.0	0.588235	...	0.0
41087	1.0	0.000000	0.636364	1.0	0.588235	...	0.0

	DayOfWeek_6	DayOfWeek_7	StoreType_a	StoreType_b
StoreType_c \				
0	0.0	0.0	0.0	0.0
1.0				
1	0.0	0.0	1.0	0.0
0.0				
2	0.0	0.0	1.0	0.0
0.0				
3	0.0	0.0	1.0	0.0
0.0				
4	0.0	0.0	1.0	0.0
0.0				
...
.				
41083	1.0	0.0	1.0	0.0
0.0				
41084	1.0	0.0	0.0	0.0

1.0				
41085	1.0	0.0	1.0	0.0
0.0				
41086	1.0	0.0	1.0	0.0
0.0				
41087	1.0	0.0	0.0	0.0
0.0				

	StoreType_d	Assortment_a	Assortment_b	Assortment_c
0	0.0	1.0	0.0	0.0
1	0.0	1.0	0.0	0.0
2	0.0	0.0	0.0	1.0
3	0.0	1.0	0.0	0.0
4	0.0	0.0	0.0	1.0
...
41083	0.0	1.0	0.0	0.0
41084	0.0	0.0	0.0	1.0
41085	0.0	0.0	0.0	1.0
41086	0.0	0.0	0.0	1.0
41087	1.0	0.0	0.0	1.0

```
[41088 rows x 24 columns],
'target_col': 'Sales',
'numeric_cols': ['Store',
'Promo',
'SchoolHoliday',
'CompetitionDistance',
'CompetitionOpen',
'Promo2',
'Day',
'Month',
'Year',
'WeekOfYear'],
'categorical_cols': ['DayOfWeek', 'StoreType', 'Assortment'],
'encoded_cols': ['DayOfWeek_1',
'DayOfWeek_2',
'DayOfWeek_3',
'DayOfWeek_4',
'DayOfWeek_5',
'DayOfWeek_6',
'DayOfWeek_7',
'StoreType_a',
'StoreType_b',
'StoreType_c',
'StoreType_d',
'Assortment_a',
'Assortment_b',
'Assortment_c']}]}
```

```
rossman_xgboost_model['train_inputs']
```

	Store	DayOfWeek	Promo	StateHoliday	Assortment	\
0	0.000000	5	1.0	0	a	
1	0.000898	5	1.0	0	a	
2	0.001795	5	1.0	0	a	
3	0.002693	5	1.0	0	c	
4	0.003591	5	1.0	0	a	
...	
1016776	0.611311	2	0.0	a	a	
1016827	0.657092	2	0.0	a	b	
1016863	0.689408	2	0.0	a	b	
1017042	0.850090	2	0.0	a	b	
1017190	0.983842	2	0.0	a	b	

	CompetitionDistance	Day	Month	Year	WeekOfYear	...	\
0	0.016482	1.0	0.545455	1.0	0.588235	...	
1	0.007252	1.0	0.545455	1.0	0.588235	...	
2	0.186050	1.0	0.545455	1.0	0.588235	...	
3	0.007911	1.0	0.545455	1.0	0.588235	...	
4	0.394119	1.0	0.545455	1.0	0.588235	...	
...	
1016776	0.001714	0.0	0.000000	0.0	0.000000	...	
1016827	0.011076	0.0	0.000000	0.0	0.000000	...	
1016863	0.010812	0.0	0.000000	0.0	0.000000	...	
1017042	0.018592	0.0	0.000000	0.0	0.000000	...	
1017190	0.009230	0.0	0.000000	0.0	0.000000	...	

	DayOfWeek_5	DayOfWeek_6	DayOfWeek_7	StoreType_a
StoreType_b \				
0	1.0	0.0	0.0	0.0
0.0				
1	1.0	0.0	0.0	1.0
0.0				
2	1.0	0.0	0.0	1.0
0.0				
3	1.0	0.0	0.0	0.0
0.0				
4	1.0	0.0	0.0	1.0
0.0				
...
.				
1016776	0.0	0.0	0.0	0.0
1.0				
1016827	0.0	0.0	0.0	0.0
1.0				
1016863	0.0	0.0	0.0	0.0
1.0				
1017042	0.0	0.0	0.0	0.0
1.0				
1017190	0.0	0.0	0.0	0.0
1.0				

	StoreType_c	StoreType_d	Assortment_a	Assortment_b
Assortment_c				
0	1.0	0.0	1.0	0.0
0.0				
1	0.0	0.0	1.0	0.0
0.0				
2	0.0	0.0	1.0	0.0
0.0				
3	1.0	0.0	0.0	0.0
1.0				
4	0.0	0.0	1.0	0.0
0.0				
...
...				
1016776	0.0	0.0	1.0	0.0
0.0				
1016827	0.0	0.0	0.0	1.0
0.0				
1016863	0.0	0.0	0.0	1.0
0.0				
1017042	0.0	0.0	0.0	1.0
0.0				
1017190	0.0	0.0	0.0	1.0
0.0				

[844392 rows x 28 columns]

```
test_preds = rossman_xgboost_model['model'].predict(X_test)
```

```
rossman_xgboost_model['model']
```

```
XGBRegressor(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=0.7, device=None,
              early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None,
              feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=0.1,
              max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=10, max_leaves=None,
              min_child_weight=None, missing=nan,
              monotone_constraints=None,
              multi_strategy=None, n_estimators=1000, n_jobs=-1,
              num_parallel_tree=None, random_state=42, ...)
```

```
test_preds
```

```
array([ 3761.39 ,  7738.725 ,  9194.781 , ...,  7101.3135,
        23356.654 ,
        7592.986 ], dtype=float32)
```

```
submission_df = pd.read_csv('csv\\sample_submission.csv')
```

```
submission_df['Sales'] = test_preds
```

```
submission_df
```

	Id	Sales
0	1	3761.389893
1	2	7738.725098
2	3	9194.781250
3	4	6687.968750
4	5	7371.160645
...
41083	41084	3162.992920
41084	41085	8253.756836
41085	41086	7101.313477
41086	41087	23356.654297
41087	41088	7592.985840

```
[41088 rows x 2 columns]
```

```
submission_df.to_csv('csv\\submission_df.csv' , index=False)
```